

# MACHINE LEARNING

## ASSIGNMENT - 4

**Name :Thandarupalli Naveen Goud**  
**UCM ID :nxt46830**  
**700# :700734683**

### Video

Link:[https://drive.google.com/file/d/1tikECPeIUwx0aZZqPNI\\_L8Gk2T-FQ8yF/view?usp=share\\_link](https://drive.google.com/file/d/1tikECPeIUwx0aZZqPNI_L8Gk2T-FQ8yF/view?usp=share_link)

Github Link: <https://github.com/nxt46830/ML-Assignment-1>

1. Apply Linear Regression to the provided dataset using underlying steps. a. Import the given “Salary\_Data.csv” b. Split the data in train\_test partitions, such that 1/3 of the data is reserved as test subset. c. Train and predict the model. d. Calculate the mean\_squared error e. Visualize both train and test data using scatter plot.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
from sklearn.decomposition import PCA
from sklearn.preprocessing import LabelEncoder, StandardScaler
import seaborn as sns
sns.set(style="white", color_codes=True)
import warnings
warnings.filterwarnings("ignore")
```

```
df=pd.read_csv("Salary_Data.csv")
df.head()
```

[1]

[2]

Visualize

	YearsExperience f.	Salary float64	
0	1.1	39343.0	
1	1.3	46205.0	
2	1.5	37731.0	
3	2.0	43525.0	
4	2.2	39891.0	

5 rows, showing 10 per page << < Page 1 of 1 > >> [Download](#)

```
X = df.iloc[:, :-1].values
Y = df.iloc[:, 1].values
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X,Y, test_size=1/3,random_state = 0)
```


```
regressor = LinearRegression()
regressor.fit(X_Train, Y_Train)

Y_Pred = regressor.predict(X_Test)
```

```
mean_squared_error(Y_Test,Y_Pred)
```

```
21026037.329511296
```

```
plt.title('Training data')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.scatter(X_Train, Y_Train)
plt.show()
```



```
plt.title('Testing data')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.scatter(X_Test, Y_Test)
plt.show()
```

[7]



2. Apply K means clustering in the dataset provided: • Remove any null values by the mean. • Use the elbow method to find a good number of clusters with the K-Means algorithm • Calculate the silhouette score for the above clustering.

```
df2=pd.read_csv("K-Mean_Dataset.csv")
df2.head()
```

	CUST_ID object	BALANCE float64	BALANCE_FREQ...	PURCHASES floa...	ONEOFF_PURC...	INSTALLMENTS...	CASH_AI
0	C10001	40.900749	0.818182	95.4	0.0	95.4	
1	C10002	3202.467416	0.909091	0.0	0.0	0.0	64
2	C10003	2495.148862	1.0	773.17	773.17	0.0	
3	C10004	1666.670542	0.636364	1499.0	1499.0	0.0	2
4	C10005	817.714335	1.0	16.0	16.0	0.0	


5 rows, showing 10 per page

```
X = df2.iloc[:,1:].values

imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
imputer = imputer.fit(X)
X = imputer.transform(X)
```

```
wcss = []
for i in range(1,11):
    kmeans = KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_state=0)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

plt.plot(range(1,11),wcss)
plt.title('the elbow method')
plt.xlabel('Number of Clusters')
plt.ylabel('Wcss')
plt.show()
```



```
from sklearn.cluster import KMeans
nclusters = 4 # this is the k in kmeans
km = KMeans(n_clusters=nclusters)
km.fit(X)
```

[11]



```
KMeans(n_clusters=4)
```

```
y_cluster_kmeans = km.predict(X)
from sklearn import metrics
score = metrics.silhouette_score(X, y_cluster_kmeans)
print('Silhouette score:', score)
```

[12]



```
Silhouette score: 0.464450664591727
```

3. Try feature scaling and then apply K-Means on the scaled features. Did that improve the Silhouette score? If Yes, can you justify why.

```
scaler = preprocessing.StandardScaler()  
scaler.fit(X)  
X_scaled_array = scaler.transform(X)  
X_scaled = pd.DataFrame(X_scaled_array)
```

[13]

```
from sklearn.cluster import KMeans  
nclusters = 4  
km = KMeans(n_clusters=nclusters)  
km.fit(X_scaled)
```

[15]

```
KMeans(n_clusters=4)
```

```
y_scaled_cluster_kmeans = km.predict(X_scaled)  
from sklearn import metrics  
score = metrics.silhouette_score(X_scaled, y_scaled_cluster_kmeans)  
print('Silhouette score after applying scaling:', score)
```

[16]

```
Silhouette score after applying scaling: 0.1976193847865969
```

The answer is “NO”, after feature scaling also it is not improving the **Silhouette Score**.